

# Building Neural Networks: A Hands-On Journey from Scratch with Python



Long Nguyen · [Follow](#)

11 min read · Nov 18, 2023



412



6






Photo by [fabio](#) on [Unsplash](#)

In this blog post, we will explore the fundamentals of neural networks, understand the intricacies of forward and backward propagation, and implement a neural network from the ground up with Python in 3 levels!

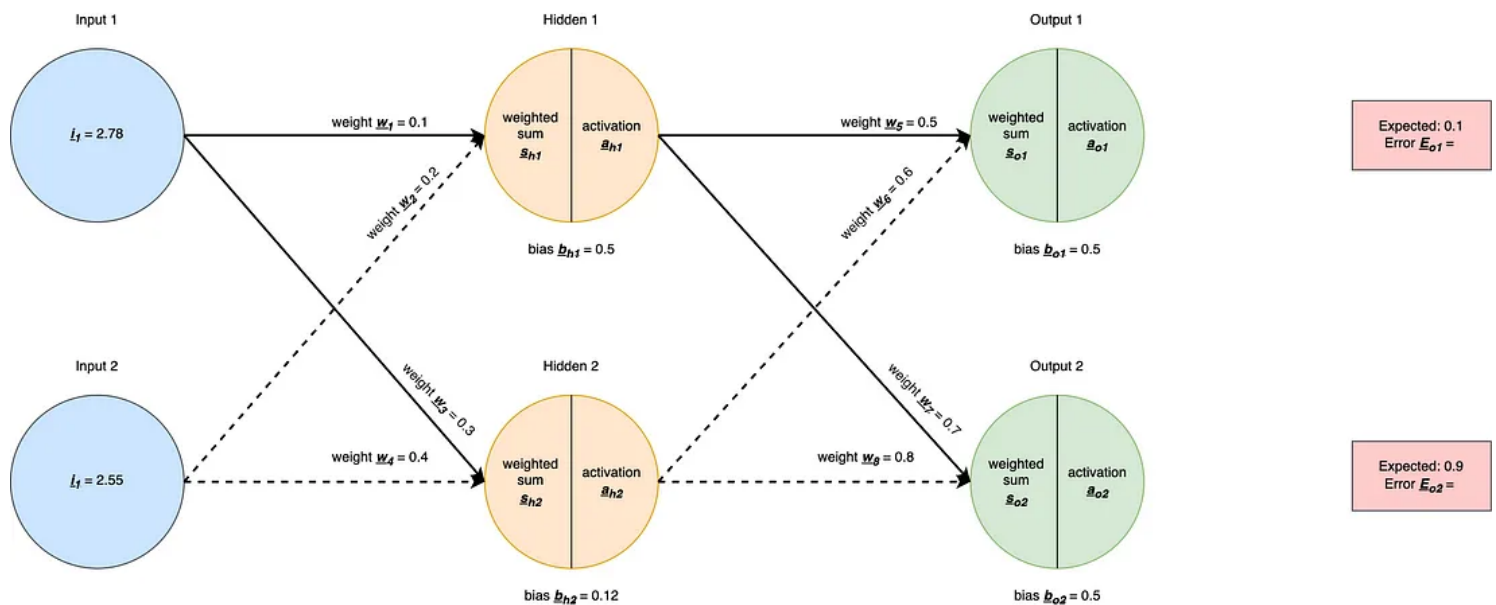
- Level 1: Without using external libraries
- Level 2: With [numpy](#)
- Level 3: With [Tensorflow](#)

If you are interested in learning about building Recurrent Neural Network from scratch as well, check out this [post](#).

# I. Forward and Backward Propagation Walkthrough

But first, let's use an example neural network and work out the mathematical calculation one neuron at a time to understand what's happening behind the scene!

Our sample neural network will consist of: 2 input neurons, 1 hidden layer with 2 neurons and an output layer with 2 neurons. Some initial weights and bias values have been provided to help with the calculation. Assume the expected output is 0.1 and 0.9:



## 1. Forward Propagation

Note: I'm using sigmoid as the activation function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

### Hidden Layer

Hidden Neuron 1:

$$WeightedSum(s_{h1}) = w_1 \times i_1 + w_2 \times i_2 + b_{h1}$$

$$WeightedSum(s_{h1}) = (2.78 \times 0.1) + (2.55 \times 0.2) + 0.5 = 0.278 + 0.51 + 0.5 = 1.288$$

$$Activation(a_{h1}) = \frac{1}{1 + e^{-s_{h1}}}$$

$$Activation(a_{h1}) = \frac{1}{1 + e^{-1.288}} \approx \frac{1}{1 + 0.275} \approx \frac{1}{1.275} \approx 0.784$$

Hidden Neuron 2:

$$WeightedSum(s_{h2}) = w_3 \times i_1 + w_4 \times i_2 + b_{h2}$$

$$WeightedSum(s_{h2}) = (2.78 \times 0.3) + (2.55 \times 0.4) + 0.12 = 0.834 + 1.02 + 0.12 = 1.974$$

$$Activation(a_{h2}) = \frac{1}{1 + e^{-s_{h2}}}$$

$$Activation(a_{h2}) = \frac{1}{1 + e^{-1.974}} \approx \frac{1}{1 + 0.138} \approx \frac{1}{1.138} \approx 0.878$$

## Output Layer

Output Neuron 1:

$$WeightedSum(s_{o1}) = w_5 \times a_{h1} + w_6 \times a_{h2} + b_{o1}$$

$$WeightedSum(s_{o1}) = (0.783 \times 0.5) + (0.878 \times 0.6) + 0.2 = 0.391 + 0.527 + 0.2 = 1.118$$

$$Activation(a_{o1}) = \frac{1}{1 + e^{-s_{o1}}}$$

$$Activation(a_{o1}) = \frac{1}{1 + e^{-1.118}} \approx \frac{1}{1 + 0.327} \approx \frac{1}{1.327} \approx 0.753$$

Output Neuron 2:

$$WeightedSum(s_{o2}) = w_7 \times a_{h1} + w_8 \times a_{h2} + b_{o2}$$

$$WeightedSum(s_{o2}) = (0.783 \times 0.7) + (0.878 \times 0.8) + 0.4 = 0.548 + 0.702 + 0.4 = 1.65$$

$$Activation(a_{o2}) = \frac{1}{1 + e^{-s_{o2}}}$$

$$Activation(a_{o2}) = \frac{1}{1 + e^{-1.65}} \approx \frac{1}{1 + 0.192} \approx \frac{1}{1.192} \approx 0.839$$

## Mean Squared Error (MSE) Calculation

Given Expected Outputs: 0.1 (for  $a_{o1}$ ) and 0.9 (for  $a_{o2}$ )

$$MSE = \frac{1}{2} \sum (Expected - Actual)^2$$

$$MSE = \frac{1}{2} ((0.1 - 0.753)^2 + (0.9 - 0.839)^2)$$

$$MSE = \frac{1}{2} (0.414 + 0.003)$$

$$MSE = 0.2085$$

So, the Mean Squared Error (MSE) is approximately 0.2085. This is a measure of the difference between the expected and actual outputs. A lower MSE

indicates a better fit of the model to the given data.

## 2. Backward Propagation

Once predictions are obtained, we need to train the network by adjusting weights and biases based on prediction errors. This is achieved through backward propagation.

Assuming we use the learning rate of 0.5

$$\alpha = 0.5$$

With backpropagation, we want to understand the sensitivity of the error function, which represents the disparity between actual and expected values, to a small adjustment (“nudge”) in a particular weight, such as  $w_5$ . I found a lot of value in revising the basics of calculus and derivatives (especially the chain rule) which has helped me grasp how backpropagation works easier. This video does a great job of explaining the intuition <https://www.youtube.com/watch?v=tIeHLnjs5U8>.

Then, the objective is to diminish the error function by reducing its gradient, thereby facilitating a ‘descent’ along the gradient — “gradient descent”.

Let’s start from the output layer and work backwards.

### Output Layer

Applying the chain rule to get the formula to calculate the change in error function with respect to a small change in weight  $w_5$

$$\frac{\partial E_{o1}}{\partial w_5} = \frac{\partial E_{o1}}{\partial a_{o1}} \cdot \frac{\partial a_{o1}}{\partial s_{o1}} \cdot \frac{\partial s_{o1}}{\partial w_5}$$

Let's work out what each component maps to.

First — We've got the error function and its derivative with respect to  $a_{o1}$

$$E = \frac{1}{2} \sum_{i=1}^n (a_{oi} - expected_i)^2$$

$$\frac{\partial E}{\partial a_{o1}} = 2 \cdot \frac{1}{2} \cdot (a_{o1} - expected_1)$$

$$\frac{\partial E_{o1}}{\partial a_{o1}} = a_{o1} - expected_1$$

Second — the derivative of the activation over the weighted sum, aka the derivative of sigmoid function

$$\frac{\partial a_{o1}}{\partial s_{o1}} = \frac{\partial}{\partial s_{o1}}(\text{Sigmoid}(s_{o1})) = a_{o1} \cdot (1 - a_{o1})$$

Lastly — the derivative of the weighted sum with respect to  $w_5$  gives you  $a_{h1}$  which is the output of the  $h1$  neuron in the previous layer

$$\frac{\partial s_{o1}}{\partial w_5} = \frac{\partial (w_5 \cdot a_{h1} + w_6 \cdot a_{h2} + b_{o1})}{\partial w_5} = a_{h1}$$

Putting them together

$$\frac{\partial E_{o1}}{\partial w_5} = (a_{o1} - \text{expected}_1) \cdot a_{o1} \cdot (1 - a_{o1}) \cdot a_{h1}$$

Usually, we can define a delta as

$$\delta_{o1} = (a_{o1} - \text{expected}_1) \cdot a_{o1} \cdot (1 - a_{o1})$$

Then the formula can be shortened to

$$\frac{\partial E_{o1}}{\partial w_5} = \delta_{o1} \cdot a_{h1}$$

This is the gradient of the error function — applying gradient descent to get a new value of weight  $w_5$  by reducing the weight it by the learning rate times gradient

$$w'_5 = w_5 - \alpha \cdot \frac{\partial E_{o1}}{\partial w_5} = w_5 - \alpha \cdot \delta_{o1} \cdot a_{h1}$$

Let's generalise the formulas for the delta of an output neuron, and the formula to update the weight in an output layer.



$$\delta_o = (a_o - \text{expected}) \cdot a_o \cdot (1 - a_o)$$

$$w_{new} = w - \alpha \cdot \delta_o \cdot a_h$$

where:

- $a_o$  is the output of the neuron,
- $\text{expected}$  is the expected output,
- $\alpha$  is the learning rate,
- $\delta_o$  is the delta of the output neuron, and
- $a_h$  is the output of the neuron in the hidden layer connected to this output neuron (aka output of the neuron in previous layer)

From this exercise, you should be able to then derive the formula for updating bias on your own — it's very similar to updating weights. Hint: the final result doesn't involve previous layer neuron's output.

Now let's apply real numbers from the example to those equations to calculate new weights  $w_5$ ,  $w_6$ ,  $w_7$ ,  $w_8$

Output Neuron 1:

$$\text{Delta } (\delta_{o1}) = (\text{Activation}(a_{o1}) - \text{Expected}(e_{o1})) \times \text{SigmoidDerivative}(a_{o1})$$

$$\text{Delta } (\delta_{o1}) : (0.753 - 0.1) \times (0.753 \times (1 - 0.753)) \approx 0.121$$

$$\text{New Weight } (w_5) : \text{Weight}(w_5) - \text{LearningRate} \times \text{Delta}(\delta_{o1}) \times \text{Activation}(a_{h1})$$

$$\text{New Weight } (w_5) : 0.5 - 0.5 \times 0.121 \times 0.784 \approx 0.453$$

$$\text{New Weight } (w_6) : \text{Weight}(w_6) - \text{LearningRate} \times \text{Delta}(\delta_{o1}) \times \text{Activation}(a_{h2})$$

$$\text{New Weight } (w_6) : 0.6 - 0.5 \times 0.121 \times 0.878 \approx 0.547$$

Output Neuron 2:

$$\text{Delta } (\delta_{o2}) = (\text{Activation}(a_{o2}) - \text{Expected}(e_{o2})) \times \text{SigmoidDerivative}(a_{o2})$$

$$\text{Delta } (\delta_{o2}) : (0.839 - 0.9) \times (0.839 \times (1 - 0.839)) \approx -0.008$$

$$\text{New Weight } (w_7) : \text{Weight}(w_7) - \text{LearningRate} \times \text{Delta}(\delta_{o2}) \times \text{Activation}(a_{h1})$$

$$\text{New Weight } (w_7) : 0.7 - 0.5 \times -0.008 \times 0.784 \approx 0.703$$

$$\text{New Weight } (w_8) : \text{Weight}(w_8) - \text{LearningRate} \times \text{Delta}(\delta_{o2}) \times \text{Activation}(a_{h2})$$

$$\text{New Weight } (w_8) : 0.8 - 0.5 \times -0.008 \times 0.878 \approx 0.804$$

## Hidden Layer

Applying the chain rule again to get the formula to calculate the change in error function with respect to a small change in weight  $w_1$

$$\frac{\partial E_{total}}{\partial w_1} = \frac{\partial E_{total}}{\partial a_{h1}} \cdot \frac{\partial a_{h1}}{\partial s_{h1}} \cdot \frac{\partial s_{h1}}{\partial w_1}$$

This formula will be a little bit more complicated as we're further away from the output, so a lot more "chaining" of functions will happen so take your time to go through this.

Looking at the first derivative — derivative of total errors with respect to  $a_{h1}$ . Because total error equals sum of  $E_{o1}$  and  $E_{o2}$ , using the sum rule we've got

$$\frac{\partial E_{total}}{\partial a_{h1}} = \frac{\partial E_{o1}}{\partial a_{h1}} + \frac{\partial E_{o2}}{\partial a_{h1}}$$

Applying the chain rule to each element

$$\frac{\partial E_{o1}}{\partial a_{h1}} = \frac{\partial E_{o1}}{\partial a_{o1}} \cdot \frac{\partial a_{o1}}{\partial s_{o1}} \cdot \frac{\partial s_{o1}}{\partial a_{h1}}$$

$$\frac{\partial E_{o2}}{\partial a_{h1}} = \frac{\partial E_{o2}}{\partial a_{o2}} \cdot \frac{\partial a_{o2}}{\partial s_{o2}} \cdot \frac{\partial s_{o2}}{\partial a_{h1}}$$

Since we calculated Delta(Eo1) and Delta (Eo2) previously

$$\delta_{o1} = \frac{\partial E_{o1}}{\partial a_{o1}} \cdot \frac{\partial a_{o1}}{\partial s_{o1}} = (a_{o1} - expected_1) \cdot a_{o1} \cdot (1 - a_{o1})$$

$$\delta_{o2} = \frac{\partial E_{o2}}{\partial a_{o2}} \cdot \frac{\partial a_{o2}}{\partial s_{o2}} = (a_{o2} - expected_2) \cdot a_{o2} \cdot (1 - a_{o2})$$

We can substitute those in

$$\frac{\partial E_{o1}}{\partial a_{h1}} = \delta_{o1} \cdot \frac{\partial s_{o1}}{\partial a_{h1}}$$

$$\frac{\partial E_{o2}}{\partial a_{h1}} = \delta_{o2} \cdot \frac{\partial s_{o2}}{\partial a_{h1}}$$

Derivative of the weighted sum with respect to previous layer neuron's output is basically just the corresponding weight

$$\frac{\partial s_{o1}}{\partial a_{h1}} = \frac{\partial (w_5 \cdot a_{h1} + w_6 \cdot a_{h2} + b_{o1})}{\partial a_{h1}} = w_5$$

$$\frac{\partial s_{o2}}{\partial a_{h1}} = \frac{\partial (w_7 \cdot a_{h1} + w_8 \cdot a_{h2} + b_{o2})}{\partial a_{h1}} = w_7$$

The derivative of total error with respect to weight *ah1* now looks like

$$\frac{\partial E_{total}}{\partial a_{h1}} = \frac{\partial E_{o1}}{\partial a_{h1}} + \frac{\partial E_{o2}}{\partial a_{h1}} = \delta_{o1} \cdot w_5 + \delta_{o2} \cdot w_7$$

Substituting this back to the initial formula

$$\frac{\partial E_{total}}{\partial w_1} = (\delta_{o1} \cdot w_5 + \delta_{o2} \cdot w_7) \cdot \frac{\partial a_{h1}}{\partial s_{h1}} \cdot \frac{\partial s_{h1}}{\partial w_1}$$

Derivative of  $ah1$  over  $sh1$  is the derivative of the sigmoid function, and the derivative of  $sh1$  over  $w1$  is the output of the previous layer neuron (which is the input layer neuron as we only have 1 hidden layer in this example)

$$\begin{aligned} \frac{\partial a_{h1}}{\partial s_{h1}} &= \frac{\partial}{\partial s_{h1}}(\text{Sigmoid}(s_{h1})) = a_{h1} \cdot (1 - a_{h1}) \\ \frac{\partial s_{h1}}{\partial w_1} &= \frac{\partial(w_1 \cdot i_1 + w_2 \cdot i_2 + b_{h1})}{\partial w_1} = i_1 \end{aligned}$$

Putting it together

$$\frac{\partial E_{total}}{\partial w_1} = (\delta_{o1} \cdot w_5 + \delta_{o2} \cdot w_7) \cdot a_{h1} \cdot (1 - a_{h1}) \cdot i_1$$

Let's group the weighted sum of the deltas in the next layer (output layer) with the sigmoid derivative and call it the  $\text{Delta}(h1)$

$$\delta_{h1} = (\delta_{o1} \cdot w_5 + \delta_{o2} \cdot w_7) \cdot a_{h1} \cdot (1 - a_{h1})$$

Rewrite the formula of the gradient of error function with respect to  $w_1$

$$\frac{\partial E_{total}}{\partial w_1} = \delta_{h1} \cdot i_1$$

Applying gradient descent and updating  $w_1$  with learning rate  $\alpha$

$$w'_1 = w_1 - \alpha \cdot \frac{\partial E_{total}}{\partial w_1} = w_1 - \alpha \cdot \delta_{h1} \cdot i_1$$

Let's generalise the formulas for the delta of a neuron in a hidden layer and the formula to update the weight in hidden layer

$$\delta_h = \left( \sum_o \delta_o \cdot w_o \right) \cdot a_h \cdot (1 - a_h)$$

$$w_{new} = w - \alpha \cdot \delta_h \cdot i$$

where:

- $\delta_o$  is the delta of an output neuron,
- $w_o$  is the weight connecting the hidden layer neuron to the output neuron
- $a_h$  is the output of the hidden layer neuron
- $a_h \cdot (1 - a_h)$  is its sigmoid derivative
- $\alpha$  is the learning rate,
- $\delta_h$  is the delta of the hidden layer neuron, and
- $i$  is the input to the hidden layer neuron (in this case where there's only 1 hidden layer, it is the input neuron, however if you have more than 1 hidden layer, it will be the output of the previous hidden layer neuron)

Now let's apply real numbers from the example to those equations to calculate new weights  $w1$ ,  $w2$ ,  $w3$ ,  $w4$

Hidden Neuron 1:

$$\text{Delta}(\delta_{h1}) (\text{Weight}(w_5) \times \text{Delta}(\delta_{o1}) + \text{Weight}(w_7) \times \text{Delta}(\delta_{o2})) \times \text{SigmoidDerivative}(a_{h1})$$

$$\text{Delta}(\delta_{h1}) : ((0.5 \times 0.121) + (0.7 \times -0.008)) \times (0.784 \times (1 - 0.784)) \approx 0.009$$

$$\text{New Weight } (w_1) : \text{Weight}(w_1) - \text{LearningRate} \times \text{Delta}(\delta_{h1}) \times \text{Input}(i_1)$$

$$\text{New Weight } (w_1) : 0.1 - 0.5 \times 0.009 \times 2.78 \approx 0.087$$

$$\text{New Weight } (w_2) : \text{Weight}(w_2) - \text{LearningRate} \times \text{Delta}(\delta_{h1}) \times \text{Input}(i_2)$$

$$\text{New Weight } (w_2) : 0.2 - 0.5 \times 0.009 \times 2.55 \approx 0.189$$

Hidden Neuron 2:

$$\text{Delta}(\delta_{h2}) : (\text{Weight}(w_6) \times \text{Delta}(\delta_{o1}) + \text{Weight}(w_8) \times \text{Delta}(\delta_{o2})) \times \text{SigmoidDerivative}(a_{h2})$$

$$\text{Delta}(\delta_{h2}) : ((0.6 \times 0.121) + (0.8 \times -0.008)) \times (0.878 \times (1 - 0.878)) \approx 0.007$$

$$\text{New Weight}(w_3) : \text{Weight}(w_3) - \text{LearningRate} \times \text{Delta}(\delta_{h2}) \times \text{Input}(i_1)$$

$$\text{New Weight}(w_3) : 0.3 - 0.5 \times 0.007 \times 2.78 \approx 0.290$$

$$\text{New Weight}(w_4) : \text{Weight}(w_4) - \text{LearningRate} \times \text{Delta}(\delta_{h2}) \times \text{Input}(i_2)$$

$$\text{New Weight}(w_4) : 0.4 - 0.5 \times 0.007 \times 2.55 \approx 0.391$$

That's it! All of our weights have been updated — and that was just 1 iteration (epoch). Imagine if we run it thousands or millions of times, the error will become smaller and smaller, hence increasing the accuracy of the network's prediction.

There were a lot of math formulas and calculations and variables so errors are quite likely to occur. If you notice something that is incorrect, please let me know!

## II. Level 1: Building A Neural Network Without Using External Libraries

Now that we've covered the math, let's dive into the first level of building a neural network: Without using external libraries (like numpy or PyTorch or Tensorflow)

First, let's define the 2 functions for the sigmoid activation function and its derivative. These 2 will be reused throughout the exercise.

```
1  from math import exp
2
3  def sigmoid(x: float) -> float:
4      return 1.0 / (1.0 + exp(-x))
5
6
7  def sigmoid_derivative(z: float) -> float:
8      return z * (1.0 - z)
```

math\_utils.py hosted with ❤ by GitHub

[view raw](#)

Now let's build a class for our Neuron:

- **Weights ( `weights` ):** Neurons receive input signals, each associated with a weight. These weights determine the importance of each input.
- **Bias ( `bias` ):** Similar to the intercept in a linear equation, the bias allows the neuron to adjust its output independently of the input.
- **Delta ( `delta` ):** This is used during the backpropagation process for adjusting weights (you see the knowledge from the walkthrough we've done earlier is coming into our code). It represents the error derivative with respect to the weighted sum.
- **Output ( `output` ):** The result of the neuron's activation function.

The sigmoid function introduces non-linearity to the model, enabling it to learn complex patterns.



```
1  from dataclasses import dataclass
2  from typing import List, Optional
3
4  from math_utils import sigmoid, sigmoid_derivative
5
6
7  @dataclass
8  class Neuron:
9      weights: List[float]
10     bias: float
11     delta: Optional[float] = 0.0
12     output: Optional[float] = 0.0
13
14     def _set_output(self, output: float) -> None:
15         self.output = output
16
17     def set_delta(self, error: float) -> None:
18         self.delta = error * sigmoid_derivative(self.output)
19
20     def weighted_sum(self, inputs: List[float]) -> float:
21         """
22         Usually results in a big number, but we tend to use a value [0, 1] for activation
23         Hence, after calculating this, we use the sigmoid function to normalize the result
24         """
25         ws = self.bias
26         for i in range(len(self.weights)):
27             ws += self.weights[i] * inputs[i]
28         return ws
29
30     def activate(self, inputs: List[float]) -> float:
31         """
32         Calculates the output of the neuron using a non-linear activation function
33         In this case we use the sigmoid function
34         """
35         output = sigmoid(self.weighted_sum(inputs))
36         self._set_output(output)
37         return output
```

neuron.py hosted with ❤ by GitHub

[view raw](#)

Neurons are then organised into layers — here's the Layer class. Layers organise neurons into meaningful groups. Neurons in the same layer share

the same input and output dimensions.

```
1  from dataclasses import dataclass
2  from typing import List, Optional
3
4  @dataclass
5  class Layer:
6      neurons: List[Neuron]
7
8      @property
9      def all_outputs(self) -> List[float]:
10         return [neuron.output for neuron in self.neurons]
11
12     def activate_neurons(self, inputs: List[float]) -> List[float]:
13         return [neuron.activate(inputs) for neuron in self.neurons]
14
15     def total_delta(self, previous_layer_neuron_idx: int) -> float:
16         return sum(
17             neuron.weights[previous_layer_neuron_idx] * neuron.delta
18             for neuron in self.neurons
19         )
```

layer.py hosted with ❤️ by GitHub

[view raw](#)

Now, the bulk of the logic is in the Network class. It represents the neural network itself and orchestrates its training and prediction processes.

Key properties:

- Hidden Layers ( `hidden_layers` ): A list containing hidden layers, each represented by the `Layer` class.
- Output Layer ( `output_layer` ): The output layer of the network, also represented by the `Layer` class.
- Learning Rate ( `learning_rate` ): A hyperparameter determining the step size at each iteration during the training process.

## Key functions:

- The `feed_forward` method conducts the forward pass, activating each neuron in sequence, starting from receiving the inputs and progressing through hidden layers to the output layer.
- The `back_propagate` method performs the backpropagation algorithm, calculating and updating the deltas of neurons in each layer. Then it calls `update_weights_for_all_layers` to update the weights after delta calculation is done.
- The `train` method trains the neural network for a specified number of epochs using the provided training set and expected outputs. The `expected` list uses one-hot encoding to indicate the expected output.

```

1  from dataclasses import dataclass
2  from typing import List, Optional
3
4
5  @dataclass
6  class Network:
7      hidden_layers: List[Layer]
8      output_layer: Layer
9      learning_rate: float
10
11  @property
12  def layers(self) -> List[Layer]:
13      return self.hidden_layers + [self.output_layer]
14
15  def feed_forward(self, inputs: List[float]) -> List[float]:
16      for layer in self.hidden_layers:
17          # update inputs as outputs of previous layers as we go
18          inputs = layer.activate_neurons(inputs)
19      return self.output_layer.activate_neurons(inputs)
20
21  def derivative_error_to_output(
22      self, actual: List[float], expected: List[float]
23  ) -> List[float]:
24      """
25      Derivative of error function with respect to the output
26      """
27      return [actual[i] - expected[i] for i in range(len(actual))]
28
29  def back_propagate(self, inputs: List[float], errors: List[float]) -> None:
30      """
31      Compute the gradient and then update the weights
32      """
33
34      # Delta of output layer = derivative of the error functions times the derivative of output
35      # We calculate deltas of output layer first
36      # So when we get to hidden layers, the output deltas are ready to be used in calculation
37      for index, neuron in enumerate(self.output_layer.neurons):
38          neuron.set_delta(errors[index])
39
40      # Calculate deltas of hidden layer
41      for layer_idx in reversed(range(len(self.hidden_layers))):
42          layer = self.hidden_layers[layer_idx]
43          next_layer = (
44              self.output_layer
45              if layer_idx == len(self.hidden_layers) - 1

```

```

43         layer_idx -= len(self.hidden_layers) - 1
44     else self.hidden_layers[layer_idx + 1]
45
46     )
47
48     for neuron_idx, neuron in enumerate(layer.neurons):
49         error_from_next_layer = next_layer.total_delta(neuron_idx)
50         neuron.set_delta(error_from_next_layer)
51
52     # Only update the weights after you've calculated the deltas
53     # If you update the weights as you move through the network, it will affect the deltas
54     self.update_weights_for_all_layers(inputs)
55
56     def update_weights_for_all_layers(self, inputs: List[float]):
57         """
58         Update weights for all layers
59         """
60         # Update weights for hidden layers
61         for layer_idx in range(len(self.hidden_layers)):
62             layer = self.hidden_layers[layer_idx]
63             previous_layer_outputs: List[float] = (
64                 inputs
65                 if layer_idx == 0
66                 else self.hidden_layers[layer_idx - 1].all_outputs
67             )
68             for neuron in layer.neurons:
69                 self.update_weights_in_a_layer(previous_layer_outputs, neuron)
70
71         # Update weights for output layer
72         for index, neuron in enumerate(self.output_layer.neurons):
73             self.update_weights_in_a_layer(self.hidden_layers[-1].all_outputs, neuron)
74
75     def update_weights_in_a_layer(
76         self, previous_layer_outputs: List[float], neuron: Neuron
77     ) -> None:
78         """
79         Update weights in all neurons in a layer
80         """
81         for idx in range(len(previous_layer_outputs)):
82             neuron.weights[idx] -= (
83                 self.learning_rate * neuron.delta * previous_layer_outputs[idx]
84             )
85             neuron.bias -= self.learning_rate * neuron.delta
86
87     def train(
88         self,
89         num_epoch: int,

```

```

90         num_outputs: int,
91         training_set: List[List[float]],
92         training_output: List[float],
93     ) -> None:
94         for epoch in range(num_epoch):
95             sum_error = 0.0
96             for idx, row in enumerate(training_set):
97                 expected = [0 for _ in range(num_outputs)]
98                 expected[training_output[idx]] = 1 # one-hot encoding
99                 actual = self.feed_forward(row)
100                 errors = self.derivative_error_to_output(actual, expected)
101                 self.back_propagate(row, errors)
102                 sum_error += self.mse(actual, training_output)
103             print(f"Mean squared error: {sum_error}")
104             print(f"epoch={epoch}")
105
106     def predict(self, inputs: List[float]) -> int:
107         outputs = self.feed_forward(inputs)
108         return outputs.index(max(outputs))
109
110     def mse(self, actual: List[float], expected: List[float]) -> float:
111         """
112         Mean Squared Error formula
113         """
114         return sum((actual[i] - expected[i]) ** 2 for i in range(len(actual))) / len(
115             actual
116         )

```

network.py hosted with ❤ by GitHub

[view raw](#)

Now it's time to try to run this code on some sample data. I've reused the data from this [tutorial](#).

This function creates a sample dataset and initialise the network with 1 hidden layer (with 2 neurons) and 1 output layer (with 2 neurons). Then, training is run for 40 epochs with learning rate of 0.5.

The neurons' weights are randomised initially and updated as training goes on.

```

1  def test_make_prediction_with_network():
2      # Test making predictions with the network
3      # Mock data is from https://machinelearningmastery.com/implement-backpropagation-algorithm-sc
4      dataset = [
5          [2.7810836, 2.550537003],
6          [1.465489372, 2.362125076],
7          [3.396561688, 4.400293529],
8          [1.38807019, 1.850220317],
9          [3.06407232, 3.005305973],
10         [7.627531214, 2.759262235],
11         [5.332441248, 2.088626775],
12         [6.922596716, 1.77106367],
13         [8.675418651, -0.242068655],
14         [7.673756466, 3.508563011],
15     ]
16     expected = [0, 0, 0, 0, 0, 1, 1, 1, 1, 1]
17     n_inputs = len(dataset[0])
18     n_outputs = len(set(expected))
19     hidden_layers = [
20         Layer(
21             neurons=[
22                 Neuron(weights=[random() for _ in range(n_inputs)], bias=random()),
23                 Neuron(weights=[random() for _ in range(n_inputs)], bias=random()),
24             ],
25         )
26     ]
27     output_layer = Layer(
28         neurons=[
29             Neuron(weights=[random() for _ in range(n_outputs)], bias=random()),
30             Neuron(weights=[random() for _ in range(n_outputs)], bias=random()),
31         ],
32     )
33     network = Network(
34         hidden_layers=hidden_layers, output_layer=output_layer, learning_rate=0.5
35     )
36     network.train(40, n_outputs, dataset, expected)
37     print(f"Hidden layer: {network.layers[0].neurons}")
38     print(f"Output layer: {network.layers[1].neurons}")
39
40     # This is just for demonstration only
41     for i in range(len(dataset)):
42         prediction = network.predict(dataset[i])
43         print("Expected=%d, Got=%d" % (expected[i], prediction))
44
45

```

```
46 if __name__ == "__main__":  
47     test_make_prediction_with_network()
```

test\_make\_prediction\_with\_network.py hosted with ❤ by GitHub

[view raw](#)

If you run the code, you should get something similar to this

```
1 Mean squared error: 3.694643563407745  
2 epoch=35  
3 Mean squared error: 3.715320273531503  
4 epoch=36  
5 Mean squared error: 3.7351519986813067  
6 epoch=37  
7 Mean squared error: 3.7541891932363085  
8 epoch=38  
9 Mean squared error: 3.7724787370507626  
10 epoch=39  
11 Hidden layer: [Neuron(weights=[-1.7955933584377923, 2.220092841757106], bias=1.8524769105456214,  
12 Output layer: [Neuron(weights=[3.8411067437736643, -0.4041378519239978], bias=-1.5094301066916858  
13 Expected=0, Got=0  
14 Expected=0, Got=0  
15 Expected=0, Got=0  
16 Expected=0, Got=0  
17 Expected=0, Got=0  
18 Expected=1, Got=1  
19 Expected=1, Got=1  
20 Expected=1, Got=1  
21 Expected=1, Got=1  
22 Expected=1, Got=1
```

sample\_terminal\_output.sh hosted with ❤ by GitHub

[view raw](#)

That concludes level 1 — building a neural network without using external libraries. As you can see, most of the math formula we derived from the initial walkthrough is used extensively in the code, so it really helps to do all of the calculation manually before you start implementing the code.



Now, the code is obviously quite lengthy and somewhat complex — let's try to simplify that by using numpy!

### **III. Level 2: Building A Neural Network With Numpy**

Since you are now familiar with the flow of the network, I'll give you all of the code at once:

```

1  from __future__ import annotations
2  from dataclasses import dataclass
3  from typing import Optional, List
4
5  import numpy as np
6
7
8  def sigmoid(x):
9      return 1 / (1 + np.exp(-x))
10
11
12 def sigmoid_derivative(z: float) -> float:
13     return z * (1.0 - z)
14
15
16 @dataclass
17 class Layer:
18     weights: np.array
19     bias: np.array
20     outputs: np.array
21     deltas: np.array
22
23
24 @dataclass
25 class Network:
26     layers: List[Layer]
27     learning_rate: Optional[int] = 0.5
28
29     @property
30     def length(self) -> int:
31         return len(self.layers)
32
33     @property
34     def outputs(self) -> np.array:
35         return self.layers[-1].outputs
36
37     @staticmethod
38     def create(
39         layers: List[int],
40     ) -> Network:
41         """
42         Create a network with random weights and biases given a list of layers
43         The "layers" is a list of the number of neurons in each layer
44         """
45         layers = [

```

```

46         Layer(
47             # layers[i] is the number of neurons in layer i (row), layers[i - 1] is the num
48             weights=np.random.rand(layers[i], layers[i - 1]),
49             bias=np.random.rand(layers[i]),
50             outputs=np.zeros(layers[i]),
51             deltas=np.zeros(layers[i]),
52         )
53     for i in range(1, len(layers))
54 ]
55
56     return Network(layers=layers)
57
58 def feed_forward(self, inputs: np.array) -> np.array:
59     for layer in self.layers:
60         # layer.outputs is a (3,1) - dimension we expect
61         # layer.weights is a (3,2), inputs is a (2,1) - multiply to get (3,1)
62         layer.outputs = sigmoid(layer.weights @ inputs + layer.bias) # == np.matmul, https
63         inputs = layer.outputs
64     return self.layers[-1].outputs
65
66 def back_propagate(self, inputs: np.array, expected: np.array) -> None:
67     for idx in reversed(range(self.length)):
68         layer = self.layers[idx]
69         if idx == len(self.layers) - 1: # if last layer (output layer)
70             layer.deltas = (layer.outputs - expected) * sigmoid_derivative(
71                 layer.outputs
72             )
73         else:
74             next_layer = self.layers[idx + 1]
75             # layer.deltas is a (3,1) - the dimension we expect
76             # next_layer.weights is a (2,3), next_layer.deltas is a (2,1)
77             # need to transpose next_layer.weights to get (3,2) then multiply by next_layer
78             layer.deltas = (
79                 next_layer.weights.T @ next_layer.deltas
80                 * sigmoid_derivative(layer.outputs)
81             ) * sigmoid_derivative(layer.outputs)
82
83     self.update_weights(inputs)
84
85 def update_weights(self, inputs: np.array) -> None:
86     for idx in range(self.length):
87         layer = self.layers[idx]
88         previous_layer_outputs = self.layers[idx - 1].outputs if idx > 0 else inputs
89         # deltas (3,) -> deltas[np.newaxis] (1, 3) -> .T (3, 1)

```

```

90     # previous_layer_outputs (2,) -> previous_layer_outputs[np.newaxis] (1, 2)
91     # (3,1) @ (1,2) = (3,2) for weights
92     layer.weights -= (
93         layer.deltas[np.newaxis].T
94         @ previous_layer_outputs[np.newaxis]
95         * self.learning_rate
96     )
97     layer.bias -= layer.deltas * self.learning_rate
98
99     def train(self, inputs: np.array, expected: np.array, epochs: int) -> None:
100         for epoch in range(epochs):
101             sum_error = 0.0
102             for idx, row in enumerate(inputs):
103                 actual = self.feed_forward(row)
104                 self.back_propagate(row, expected[idx])
105                 sum_error += self.mse(actual, expected[idx])
106             print(f"Mean squared error: {sum_error}")
107             print(f"epoch={epoch}")
108
109         def mse(self, actual: np.array, expected: np.array) -> float:
110             return np.power(actual - expected, 2).mean()
111
112         def predict(self, inputs: np.array) -> int:
113             outputs = self.feed_forward(inputs)
114             return np.where(outputs == outputs.max())[0][0]

```

nn\_with\_numpy.py hosted with ❤ by GitHub

[view raw](#)

Using numpy helps us shorten the code a little bit — you can imagine that it's doing “bulk” calculation by utilising matrices instead of looping one neuron and one layer at a time as our previous implementation.

However, you'd need to have a pretty good mental model of the dimensions of the matrices in each step in order to understand and write the correct calculation, which can be a bit challenging. I've put comments in the code about the dimensions expected for most of the calculations (based on the test case).

The `Layer` class is a data class that encapsulates the parameters and attributes associated with a layer in the neural network. We don't need a `Neuron` class anymore since it will just be an element in the numpy array/matrix.

I've added a static method `create` which creates a network with random weights and biases based on the specified number of neurons in each layer. The rest of the functions are the same, except the calculation is done with matrix multiplications instead of manually multiplying each neuron's data.

There are obviously many ways of implementing this — one might simplify this further and remove the `Layer` class completely and represents the whole network with nested arrays. However I find that approach a bit hard to wrap my head around with all of the multiple dimensions so I went with this approach for now.

Let's try running the code with the same dataset

```

1
2 def test_make_prediction_with_network():
3     # Test making predictions with the network
4     # Mock data is from https://machinelearningmastery.com/implement-backpropagation-algorithm-s
5     dataset = np.array(
6         [
7             [2.7810836, 2.550537003],
8             [1.465489372, 2.362125076],
9             [3.396561688, 4.400293529],
10            [1.38807019, 1.850220317],
11            [3.06407232, 3.005305973],
12            [7.627531214, 2.759262235],
13            [5.332441248, 2.088626775],
14            [6.922596716, 1.77106367],
15            [8.675418651, -0.242068655],
16            [7.673756466, 3.508563011],
17        ]
18    )
19    expected = np.array(
20        [
21            [1, 0],
22            [1, 0],
23            [1, 0],
24            [1, 0],
25            [1, 0],
26            [0, 1],
27            [0, 1],
28            [0, 1],
29            [0, 1],
30            [0, 1],
31        ]
32    )
33    # 2 input neurons, 3 hidden neurons, 2 output neurons
34    network = Network.create([len(dataset[0]), 3, len(expected[0])])
35    network.train(dataset, expected, 40)
36    for i in range(len(dataset)):
37        prediction = network.predict(dataset[i])
38        print(
39            f"{i} - Expected={np.where(expected[i] == expected[i].max())[0][0]}, Got={prediction}
40        )
41
42
43 if __name__ == "__main__":
44     test_make_prediction_with_network()

```

[test\\_make\\_prediction\\_with\\_network\\_numpy.py](#) hosted with ❤ by GitHub[view raw](#)

The dimensions of input and expected output are obviously slightly changed to fit with numpy n-D arrays, but the data stay the same.

This is a sample output

```
1 Mean squared error: 1.1978125477947856
2 epoch=36
3 Mean squared error: 1.1020105642439684
4 epoch=37
5 Mean squared error: 1.013486082749448
6 epoch=38
7 Mean squared error: 0.9322033466403574
8 epoch=39
9 0 - Expected=0, Got=0
10 1 - Expected=0, Got=0
11 2 - Expected=0, Got=0
12 3 - Expected=0, Got=0
13 4 - Expected=0, Got=0
14 5 - Expected=1, Got=1
15 6 - Expected=1, Got=1
16 7 - Expected=1, Got=1
17 8 - Expected=1, Got=1
18 9 - Expected=1, Got=1
```

[numpy-sample-output.sh](#) hosted with ❤ by GitHub[view raw](#)

### III. Level 3: Building A Neural Network With Tensorflow

In this level, we've transitioned from a detailed, 200-line implementation of a neural network to just a few concise lines using TensorFlow. The power of TensorFlow allows us to express neural network architectures with ease.

However, I won't go into details about this code, and Tensorflow in general as our aim in this blog is not to delve into the intricacies of TensorFlow but to comprehend the fundamental workings of a neural network. TensorFlow abstracts away many of the underlying details, making it an efficient tool for practical applications but potentially not the best way to learn.

This is merely to demonstrate that neural network is complex and to fully understand it, it's recommended to attempt to build one from scratch. Starting from the basics lays a solid foundation, enabling a deeper understanding of the complexities involved. While libraries like TensorFlow offer convenience, diving into their usage without a fundamental understanding of neural networks can hinder comprehensive learning.



```

1  import numpy as np
2  import tensorflow as tf
3  from tensorflow import keras
4  from tensorflow.keras import layers
5
6
7  def build_model() -> tf.keras.Sequential:
8      model = tf.keras.Sequential(
9          [
10             layers.Dense(units=3, activation="sigmoid", input_shape=(2,)),
11             layers.Dense(units=2),
12         ]
13     )
14     model.summary()
15     loss_fn = keras.losses.MeanSquaredError()
16     model.compile(optimizer="adam", loss=loss_fn, metrics=["accuracy"])
17     return model
18
19
20 if __name__ == "__main__":
21     dataset = np.array(
22         [
23             [2.7810836, 2.550537003],
24             [1.465489372, 2.362125076],
25             [3.396561688, 4.400293529],
26             [1.38807019, 1.850220317],
27             [3.06407232, 3.005305973],
28             [7.627531214, 2.759262235],
29             [5.332441248, 2.088626775],
30             [6.922596716, 1.77106367],
31             [8.675418651, -0.242068655],
32             [7.673756466, 3.508563011],
33         ]
34     )
35     expected = np.array(
36         [
37             [1, 0],
38             [1, 0],
39             [1, 0],
40             [1, 0],
41             [1, 0],
42             [0, 1],
43             [0, 1],
44             [0, 1],
45             [0, 1]

```

```
45         [0, 1],
46     ]
47
48 )
49 model = build_model()
50 # Convert the data to TensorFlow format
51 dataset_tf = tf.constant(dataset, dtype=tf.float32)
52 expected_tf = tf.constant(expected, dtype=tf.float32)
53
54 # Train the model
55 model.fit(dataset_tf, expected_tf, epochs=200)
56 predictions = model.predict(dataset_tf)
57 print(predictions)
58
59 # Convert continuous predictions to class labels (0 or 1)
60 class_predictions = np.argmax(predictions, axis=1)
61 print(class_predictions)
62
63 # Print the comparison
64 for i, (expected_row, prediction) in enumerate(zip(expected, class_predictions)):
65     print(f"{i} - Expected={expected_row.argmax()}, Got={prediction}")
```

nn\_with\_tf.py hosted with ❤️ by GitHub

[view raw](#)

## Sample output

```
1 Epoch 196/200
2 1/1 [=====] - 0s 784us/step - loss: 0.0737 - accuracy: 1.0000
3 Epoch 197/200
4 1/1 [=====] - 0s 691us/step - loss: 0.0733 - accuracy: 1.0000
5 Epoch 198/200
6 1/1 [=====] - 0s 643us/step - loss: 0.0729 - accuracy: 1.0000
7 Epoch 199/200
8 1/1 [=====] - 0s 629us/step - loss: 0.0725 - accuracy: 1.0000
9 Epoch 200/200
10 1/1 [=====] - 0s 640us/step - loss: 0.0721 - accuracy: 1.0000
11 1/1 [=====] - 0s 27ms/step
12 [[0.66456985 0.39463678]
13  [0.75454247 0.1885925 ]
14  [0.92614865 0.29184788]
15  [0.6771406  0.21432315]
16  [0.7149838  0.39005262]
17  [0.17372087 0.8059198 ]
18  [0.2435018  0.72199327]
19  [0.19464082 0.74787194]
20  [0.2519498  0.5949259 ]
21  [0.17731054 0.8268534 ]]
22 [0 0 0 0 1 1 1 1 1]
23 0 - Expected=0, Got=0
24 1 - Expected=0, Got=0
25 2 - Expected=0, Got=0
26 3 - Expected=0, Got=0
27 4 - Expected=0, Got=0
28 5 - Expected=1, Got=1
29 6 - Expected=1, Got=1
30 7 - Expected=1, Got=1
31 8 - Expected=1, Got=1
32 9 - Expected=1, Got=1
```

sample\_terminal\_output\_tf.sh hosted with ❤ by GitHub

[view raw](#)

In this blog journey, we took a dive into the behind the scene of neural networks, starting from the basic walkthrough with math calculation and then moving into code implementation with Python. We built these

Open in app ↗



Search

Write



As we wrap up, I invite you to try this out for yourself. Coding is a journey of discovery, and building a neural network from scratch is like a backstage tour. So, grab your coding gear, start tinkering, and enjoy the adventure of learning. Happy coding!

### References and resources:

- [Gradient descent, how neural networks learn](#)
- [What is backpropagation really doing?](#)
- [The essence of calculus](#)
- [Backpropagation calculus](#)
- [A Step by Step Backpropagation Example](#)
- [How to Code a Neural Network with Backpropagation In Python \(from scratch\)](#)
- [Difference between numpy dot\(\) and Python 3.5+ matrix multiplication](#)
- [CHAPTER 2 — How the backpropagation algorithm works](#)

Neural Networks

Backpropagation

Numpy

Artificial Neural Network

Deep Learning



**Written by Long Nguyen**

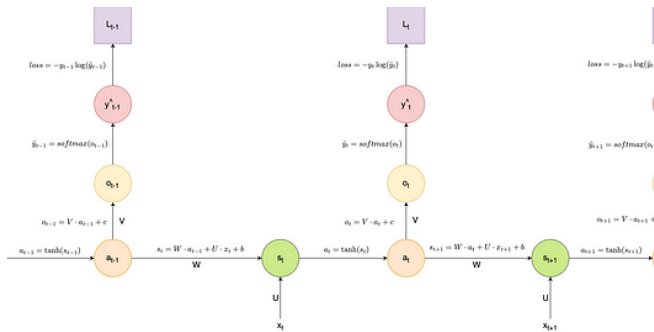
202 Followers

Software Engineer @ Atlassian



---

**More from Long Nguyen**



Long Nguyen

## Building a Recurrent Neural Network From Scratch

In this blog post, we will explore Recurrent Neural Networks (RNNs) and the...


14 min read · Jan 28, 2024



77



...



Sinem Günel in The Startup

## I Constantly Use MrBeast's Strategies to Grow My Business—...

3 psychological biases you won't be able to unsee anymore

🌟 · 11 min read · Feb 13, 2024



5.1K



86



...



Desiree Peralta in The Startup

## One of The Most Profitable 1-Person Business Anyone Can Star...

Regardless of your knowledge, experience, and social media presence.

🌟 · 7 min read · Feb 2, 2024



2.6K



51



...



Long Nguyen in The Startup

## Large Language Models & Transformer Architecture: The...

Explore LLMs and the Transformer architecture: from tokenisation, embeddings...

15 min read · Jul 25, 2023



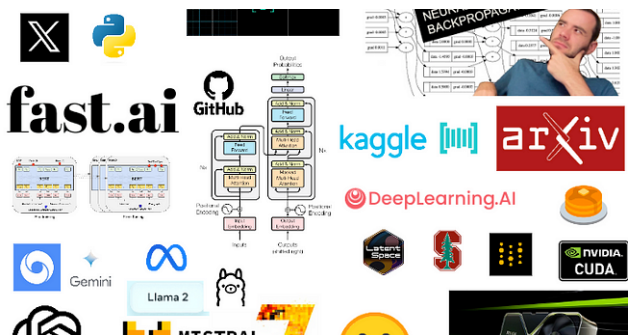
16



...

[See all from Long Nguyen](#)

## Recommended from Medium



Benedict Neo in bitgrit Data Science Publication

### Roadmap to Learn AI in 2024

A free curriculum for hackers and programmers to learn AI

11 min read · 3 days ago



2.2K



26



Koushik

### Understanding Convolutional Neural Networks (CNNs) in Depth

Convolutional Neural Networks skillfully capturing and extracting patterns from data,...

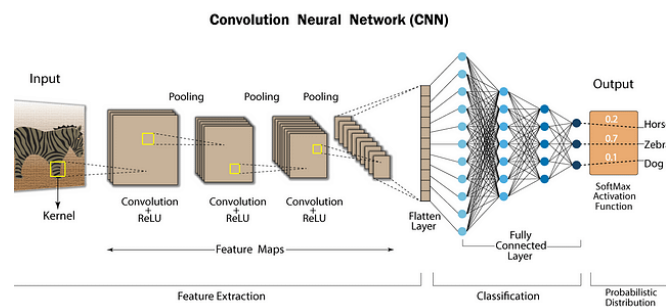
12 min read · Nov 28, 2023




583



5



## Lists




### Natural Language Processing

1224 stories · 706 saves



### Practical Guides to Machine Learning

10 stories · 1096 saves



## data science and AI

40 stories · 83 saves



## Staff Picks

587 stories · 776 saves



Fareed Khan in Level Up Coding

## Solving Transformer by Hand: A Step-by-Step Math Example

Performing numerous matrix multiplications to solve the encoder and decoder parts of th...

13 min read · Dec 18, 2023



1.7K



25



...



Jason Roell

## Ultimate Python Cheat Sheet: Practical Python For Everyday...

This Cheat Sheet was born out of necessity. Recently, I was tasked with diving into a new...

33 min read · Jan 30, 2024




2.1K



27



...




Asad iqbal

## Demystifying Machine Learning: A Guided Tour of the Top 10...

1. Linear Regression

19 min read · Sep 16, 2023



Luís Fernando Torres in LatinXinAI

## Convolutional Neural Network From Scratch

The most effective way of working with image data

21 min read · Oct 16, 2023





---

See more recommendations